

Machine Learning Lab 1

Akhilendra Pratap 211112438

```
import numpy as np
import pandas as pd
import matplotlib.pyplot as plt
import math
import os
```

```
/tmp/ipykernel_5139/4209716058.py:2: DeprecationWarning:
Pyarrow will become a required dependency of pandas in the next major
release of pandas (pandas 3.0),
(to allow more performant data types, such as the Arrow string type,
and better interoperability with other libraries)
but was not found to be installed on your system.
If this would cause problems for you,
please provide us feedback at
https://github.com/pandas-dev/pandas/issues/54466
```

```
import pandas as pd

train_dirs=[]
test_dirs=[]
for dir in os.listdir("./"):
    if(dir.find("5-fold")!=-1):
        train_dirs.append("./"+dir+"/train/")
        test_dirs.append("./"+dir+"/test/")

def cal_header_val(file_path):
    with open(file_path, "r") as file:
        lines=file.readlines()
    return lines.index('@data\n')+1

headers=[]
for dir in train_dirs:
    file_path=dir+os.listdir (dir)[0]
    headers.append(cal_header_val(file_path))

headers=np.array(headers)
```

Linear Regression

```
def LinearRegression(train_file, test_file, header):
    train_df = pd.read_csv(train_file, header=header, delimiter=",")
    test_df = pd.read_csv(test_file, header=header, delimiter=",")

    X_train = train_df.iloc[:, :-1].values
    y_train = train_df.iloc[:, -1].values
```

```

X_test = test_df.iloc[:, :-1].values
y_test = test_df.iloc[:, -1].values

from sklearn.linear_model import LinearRegression
regressor = LinearRegression()

regressor.fit(X_train, y_train)

y_pred = regressor.predict(X_test)

from sklearn.metrics import mean_squared_error,
mean_absolute_error, r2_score
mse = mean_squared_error(y_test, y_pred)
mae = mean_absolute_error(y_test, y_pred)
r2score = r2_score(y_test, y_pred)
return np.array([mse, mae, r2score])

for train_dir, test_dir, header in zip(train_dirs, test_dirs,
headers):
    train_files=os.listdir(train_dir)
    test_files=os.listdir(test_dir)
    val=np.zeros(3)
    for train, test in zip(train_files, test_files):
        val+=LinearRegression(train_dir+train, test_dir+test, header)
    print(train_dir)
    val/=len(train_files)
    val[0]=math.sqrt(val[0])
    val=pd.DataFrame(val, index=["RMSE", "MSE", "R2"],
columns=["Values"])
    print(val)
    print("-----\n")

```

```

/usr/lib/python3/dist-packages/scipy/__init__.py:146: UserWarning: A
NumPy version >=1.17.3 and <1.25.0 is required for this version of
SciPy (detected version 1.26.3

```

```

    warnings.warn(f"A NumPy version >={np_minversion} and
<{np_maxversion}")

```

```

./ele-1-5-fold/train/

```

```

      Values
RMSE  633.808233
MSE    419.216579
R2       0.697299
-----

```

```

./diabetes-5-fold/train/

```

```

      Values
RMSE   0.589029
MSE    0.465676
R2     0.144824

```

```

-----
./quake-5-fold/train/
      Values
RMSE   0.188958
MSE    0.148494
R2     0.003954
-----

./laser-5-fold/train/
      Values
RMSE   23.022278
MSE    15.511976
R2     0.752009
-----

./plastic-5-fold/train/
      Values
RMSE   1.529883
MSE    1.231589
R2     0.798854
-----

```

Polynomial Regression of degree 2 and 3

```

def PolynomialRegression(train_file, test_file, header, degree):
    train_df = pd.read_csv(train_file, header=header, delimiter=",")
    test_df = pd.read_csv(test_file, header=header, delimiter=",")

    X_train = train_df.iloc[:, :-1].values
    y_train = train_df.iloc[:, -1].values

    X_test = test_df.iloc[:, :-1].values
    y_test = test_df.iloc[:, -1].values

    from sklearn.linear_model import LinearRegression
    from sklearn.preprocessing import PolynomialFeatures

    poly_reg=PolynomialFeatures(degree=degree)
    X_poly=poly_reg.fit_transform(X_train)

    regressor = LinearRegression()
    regressor.fit(X_poly, y_train)

    y_pred = regressor.predict(poly_reg.transform(X_test))

    from sklearn.metrics import mean_squared_error,
    mean_absolute_error, r2_score

```

```

mse = mean_squared_error(y_test, y_pred)
mae = mean_absolute_error(y_test, y_pred)
r2score = r2_score(y_test, y_pred)
return np.array([mse, mae, r2score])

```

Polynomial Regression of degree 2

```

for train_dir, test_dir, header in zip(train_dirs, test_dirs,
headers):
    train_files = os.listdir(train_dir)
    test_files = os.listdir(test_dir)

    val = np.zeros(3)
    for train, test in zip(train_files, test_files):
        val += PolynomialRegression(train_dir + train, test_dir +
test, header, 2)
    print(train_dir)
    val /= len(train_files)
    val[0] = math.sqrt(val[0])
    val = pd.DataFrame(val, index=["RMSE", "MSE", "R2"],
columns=["Values"])
    print(val)
    print("-----\n")

```

```

./ele-1-5-fold/train/
      Values
RMSE  593.068697
MSE   404.025764
R2     0.732678
-----

```

```

./diabetes-5-fold/train/
      Values
RMSE  0.507230
MSE   0.414338
R2    0.355809
-----

```

```

./quake-5-fold/train/
      Values
RMSE  0.188900
MSE   0.148034
R2    0.004362
-----

```

```

./laser-5-fold/train/
      Values
RMSE  10.382359
MSE   6.491844
R2    0.949910

```

```

-----
./plastic-5-fold/train/
      Values
RMSE  1.524322
MSE    1.222752
R2     0.800348
-----

```

Polynomial Regression of degree 3

```

for train_dir, test_dir, header in zip(train_dirs, test_dirs,
headers):
    train_files = os.listdir(train_dir)
    test_files = os.listdir(test_dir)

    val = np.zeros(3)
    for train, test in zip(train_files, test_files):
        val += PolynomialRegression(train_dir + train, test_dir +
test, header, 3)
    print(train_dir)
    val /= len(train_files)
    val[0] = math.sqrt(val[0])
    val = pd.DataFrame(val, index=["RMSE", "MSE", "R2"],
columns=["Values"])
    print(val)
    print("-----\n")

```

```

./ele-1-5-fold/train/
      Values
RMSE  584.772891
MSE   397.998001
R2     0.739070
-----

```

```

./diabetes-5-fold/train/
      Values
RMSE  0.479272
MSE   0.364019
R2    0.463297
-----

```

```

./quake-5-fold/train/
      Values
RMSE  0.188188
MSE   0.148102
R2    0.011537
-----

```

```
./laser-5-fold/train/
```

```
      Values  
RMSE  5.376591  
MSE   2.819652  
R2    0.986938  
-----
```

```
./plastic-5-fold/train/
```

```
      Values  
RMSE  1.468539  
MSE   1.161601  
R2    0.814599  
-----
```

Regularization in Linear Regression

```
def Regularization(train_file, test_file, header):  
    train_df = pd.read_csv(train_file, header=header, delimiter=",")  
    test_df = pd.read_csv(test_file, header=header, delimiter=",")  
  
    X_train = train_df.iloc[:, :-1].values  
    y_train = train_df.iloc[:, -1].values  
  
    X_test = test_df.iloc[:, :-1].values  
    y_test = test_df.iloc[:, -1].values  
  
    from sklearn.linear_model import Ridge  
    alphas = np.array([2**i for i in range(-18, 51, 2)])  
  
    best_mse, best_alpha=float('inf'), None  
  
    from sklearn.metrics import mean_squared_error  
    for alpha in alphas:  
        regressor=Ridge(alpha=alpha)  
        regressor.fit(X_train, y_train)  
        y_pred=regressor.predict(X_test)  
        mse=mean_squared_error(y_test, y_pred)  
  
        if mse<best_mse:  
            best_mse, best_alpha=mse, alpha  
    return np.array([best_mse, best_alpha])  
  
for train_dir, test_dir, header in zip(train_dirs, test_dirs,  
headers):  
    train_files = os.listdir(train_dir)  
    test_files = os.listdir(test_dir)  
  
    val = np.zeros(2)  
    for train, test in zip(train_files, test_files):
```

```

        val = Regularization(train_dir + train, test_dir + test,
header)
        print(train_dir)
        val[0]=math.sqrt(val[0])
        val = pd.DataFrame(val, index=["Best RMSE", "Best Alpha"],
columns=["Values"])
        print(val)
        print("-----\n")

```

```

./ele-1-5-fold/train/
          Values
Best RMSE   575.377949
Best Alpha    0.000004
-----

```

```

./diabetes-5-fold/train/
          Values
Best RMSE    0.634371
Best Alpha    0.000004
-----

```

```

./quake-5-fold/train/
          Values
Best RMSE    0.178702
Best Alpha    0.000004
-----

```

```

./laser-5-fold/train/
          Values
Best RMSE    24.206672
Best Alpha  262144.000000
-----

```

```

./plastic-5-fold/train/
          Values
Best RMSE    1.510579
Best Alpha    0.000004
-----

```

Ridge on Linear Reg for Best MAE

```

def Regularization(train_file, test_file, header):
    train_df = pd.read_csv(train_file, header=header, delimiter=",")
    test_df = pd.read_csv(test_file, header=header, delimiter=",")

    X_train = train_df.iloc[:, :-1].values
    y_train = train_df.iloc[:, -1].values

    X_test = test_df.iloc[:, :-1].values

```

```

y_test = test_df.iloc[:, -1].values

from sklearn.linear_model import Ridge
alphas = np.array([2**i for i in range(-18, 51, 2)])

best_mse, best_alpha=float('inf'), None

from sklearn.metrics import mean_absolute_error
for alpha in alphas:
    regressor=Ridge(alpha=alpha)
    regressor.fit(X_train, y_train)
    y_pred=regressor.predict(X_test)
    mse=mean_absolute_error(y_test, y_pred)

    if mse<best_mse:
        best_mse, best_alpha=mse, alpha
return np.array([best_mse, best_alpha])

for train_dir, test_dir, header in zip(train_dirs, test_dirs,
headers):
    train_files = os.listdir(train_dir)
    test_files = os.listdir(test_dir)

    val = np.zeros(2)
    for train, test in zip(train_files, test_files):
        val = Regularization(train_dir + train, test_dir + test,
header)
    print(train_dir)
    val = pd.DataFrame(val, index=["Best MAE", "Best Alpha"],
columns=["Values"])
    print(val)
    print("-----\n")

```

```

./ele-1-5-fold/train/
          Values
Best MAE      397.714181
Best Alpha    65536.000000
-----

```

```

./diabetes-5-fold/train/
          Values
Best MAE      0.47644
Best Alpha     4.00000
-----

```

```

./quake-5-fold/train/
          Values
Best MAE      0.139492
Best Alpha    65536.000000
-----

```



```

./laser-5-fold/train/
                Values
Best MAE        15.894284
Best Alpha  262144.000000
-----

./plastic-5-fold/train/
                Values
Best MAE        1.234517
Best Alpha   64.000000
-----

```

Ridge for Linear Reg, Best R2 score

```

def Regularization(train_file, test_file, header):
    train_df = pd.read_csv(train_file, header=header, delimiter=",")
    test_df = pd.read_csv(test_file, header=header, delimiter=",")

    X_train = train_df.iloc[:, :-1].values
    y_train = train_df.iloc[:, -1].values

    X_test = test_df.iloc[:, :-1].values
    y_test = test_df.iloc[:, -1].values

    from sklearn.linear_model import Ridge
    alphas = np.array([2**i for i in range(-18, 51, 2)])

    best_mse, best_alpha = float('inf'), None

    from sklearn.metrics import r2_score
    for alpha in alphas:
        regressor = Ridge(alpha=alpha)
        regressor.fit(X_train, y_train)
        y_pred = regressor.predict(X_test)
        mse = r2_score(y_test, y_pred)

        if mse < best_mse:
            best_mse, best_alpha = mse, alpha
    return np.array([best_mse, best_alpha])

for train_dir, test_dir, header in zip(train_dirs, test_dirs,
headers):
    train_files = os.listdir(train_dir)
    test_files = os.listdir(test_dir)

    val = np.zeros(2)
    for train, test in zip(train_files, test_files):
        val = Regularization(train_dir + train, test_dir + test,

```

```

header)
    print(train_dir)
    val = pd.DataFrame(val, index=["Best R2 Score", "Best Alpha"],
columns=["Values"])
    print(val)
    print("-----\n")

./ele-1-5-fold/train/
          Values
Best R2 Score  0.738707
Best Alpha    0.000004
-----

./diabetes-5-fold/train/
          Values
Best R2 Score  0.319099
Best Alpha    0.000004
-----

./quake-5-fold/train/
          Values
Best R2 Score  0.005949
Best Alpha    0.000004
-----

./laser-5-fold/train/
          Values
Best R2 Score  0.662978
Best Alpha    262144.000000
-----

./plastic-5-fold/train/
          Values
Best R2 Score  0.790441
Best Alpha    0.000004
-----

import warnings
warnings.filterwarnings('ignore')

```

Ridge for All using MSE

```

def PolynomialRidge(train_file, test_file, header, degree):
    train_df = pd.read_csv(train_file, header=header, delimiter=",")
    test_df = pd.read_csv(test_file, header=header, delimiter=",")

    X_train = train_df.iloc[:, :-1].values
    y_train = train_df.iloc[:, -1].values

```

```

X_test = test_df.iloc[:, :-1].values
y_test = test_df.iloc[:, -1].values

from sklearn.linear_model import Ridge
from sklearn.preprocessing import PolynomialFeatures

poly_reg=PolynomialFeatures(degree=degree)
X_poly=poly_reg.fit_transform(X_train)

alphas=np.array([2**i for i in range(-18, 30)])

best_alpha, best_mse=None, float('inf')
for alpha in alphas:
    regressor = Ridge(alpha=alpha)
    regressor.fit(X_poly, y_train)
    y_pred = regressor.predict(poly_reg.transform(X_test))

    from sklearn.metrics import mean_squared_error,
mean_absolute_error, r2_score
    mse = mean_squared_error(y_test, y_pred)
    if(mse<best_mse):
        best_mse, best_alpha=mse, alpha

    return np.array([best_mse, best_alpha])

for train_dir, test_dir, header in zip(train_dirs, test_dirs,
headers):
    train_files = os.listdir(train_dir)
    test_files = os.listdir(test_dir)

    val = np.zeros(2)
    for train, test in zip(train_files, test_files):
        val += PolynomialRidge(train_dir + train, test_dir + test,
header, 2)
    print(train_dir)
    val /= len(train_files)
    val[0] = math.sqrt(val[0])
    val = pd.DataFrame(val, index=["RMSE", "Alpha"],
columns=["Values"])
    print(val)
    print("-----\n")

./ele-1-5-fold/train/
      Values
RMSE      592.073943
Alpha  56524.800002
-----

./diabetes-5-fold/train/
      Values

```

```
RMSE      0.505328
Alpha    419533.200002
-----
```

```
./quake-5-fold/train/
      Values
```

```
RMSE    1.887830e-01
Alpha   1.385431e+07
-----
```

```
./laser-5-fold/train/
      Values
```

```
RMSE     10.373138
Alpha    217.600002
-----
```

```
./plastic-5-fold/train/
      Values
```

```
RMSE      1.523200
Alpha    26252.800002
-----
```

```
for train_dir, test_dir, header in zip(train_dirs, test_dirs,
headers):
    train_files = os.listdir(train_dir)
    test_files = os.listdir(test_dir)

    val = np.zeros(2)
    for train, test in zip(train_files, test_files):
        val += PolynomialRidge(train_dir + train, test_dir + test,
header, 3)
    print(train_dir)
    val /= len(train_files)
    val[0] = math.sqrt(val[0])
    val = pd.DataFrame(val, index=["RMSE", "Alpha"],
columns=["Values"])
    print(val)
    print("-----\n")
```

```
./ele-1-5-fold/train/
      Values
```

```
RMSE     5.845551e+02
Alpha    1.074135e+08
-----
```

```
./diabetes-5-fold/train/
      Values
```

```
RMSE      0.455948
Alpha    218.403126
```

```

-----
./quake-5-fold/train/
      Values
RMSE   1.879962e-01
Alpha  1.109393e+08
-----

```

```

./laser-5-fold/train/
      Values
RMSE   5.236308
Alpha  52.800002
-----

```

```

./plastic-5-fold/train/
      Values
RMSE   1.468035
Alpha  25.625002
-----

```

Ridge for all using MAE

```

def PolynomialRidge(train_file, test_file, header, degree):
    train_df = pd.read_csv(train_file, header=header, delimiter=",")
    test_df = pd.read_csv(test_file, header=header, delimiter=",")

    X_train = train_df.iloc[:, :-1].values
    y_train = train_df.iloc[:, -1].values

    X_test = test_df.iloc[:, :-1].values
    y_test = test_df.iloc[:, -1].values

    from sklearn.linear_model import Ridge
    from sklearn.preprocessing import PolynomialFeatures

    poly_reg=PolynomialFeatures(degree=degree)
    X_poly=poly_reg.fit_transform(X_train)

    alphas=np.array([2**i for i in range(-18, 30)])

    best_alpha, best_mse=None, float('inf')
    for alpha in alphas:
        regressor = Ridge(alpha=alpha)
        regressor.fit(X_poly, y_train)
        y_pred = regressor.predict(poly_reg.transform(X_test))

        from sklearn.metrics import mean_squared_error,
        mean_absolute_error, r2_score
        mse = mean_absolute_error(y_test, y_pred)

```

```

        if(mse<best_mse):
            best_mse, best_alpha=mse, alpha

    return np.array([best_mse, best_alpha])

for train_dir, test_dir, header in zip(train_dirs, test_dirs,
headers):
    train_files = os.listdir(train_dir)
    test_files = os.listdir(test_dir)

    val = np.zeros(2)
    for train, test in zip(train_files, test_files):
        val += PolynomialRidge(train_dir + train, test_dir + test,
header, 3)
    print(train_dir)
    val /= len(train_files)
    val = pd.DataFrame(val, index=["MAE", "Alpha"],
columns=["Values"])
    print(val)
    print("-----\n")

```

```
./ele-1-5-fold/train/
```

```

          Values
MAE      397.745854
Alpha  420352.000002
-----

```

```
./diabetes-5-fold/train/
```

```

          Values
MAE       0.353695
Alpha  52441.600002
-----

```

```
./quake-5-fold/train/
```

```

          Values
MAE    1.480013e-01
Alpha  3.460301e+06
-----

```

```
./laser-5-fold/train/
```

```

          Values
MAE       2.816814
Alpha  12.800003
-----

```

```
./plastic-5-fold/train/
```

```

          Values
MAE       1.160367
Alpha    0.262501

```

```

-----

for train_dir, test_dir, header in zip(train_dirs, test_dirs,
headers):
    train_files = os.listdir(train_dir)
    test_files = os.listdir(test_dir)

    val = np.zeros(2)
    for train, test in zip(train_files, test_files):
        val += PolynomialRidge(train_dir + train, test_dir + test,
header, 2)
    print(train_dir)
    val /= len(train_files)
    val = pd.DataFrame(val, index=["MAE", "Alpha"],
columns=["Values"])
    print(val)
    print("-----\n")

```

```

./ele-1-5-fold/train/
      Values
MAE      401.655189
Alpha  59801.600002
-----

```

```

./diabetes-5-fold/train/
      Values
MAE      0.401606
Alpha  419841.600002
-----

```

```

./quake-5-fold/train/
      Values
MAE      1.479610e-01
Alpha  6.710886e+06
-----

```

```

./laser-5-fold/train/
      Values
MAE      6.439135
Alpha  1356.800000
-----

```

```

./plastic-5-fold/train/
      Values
MAE      1.219512
Alpha  26368.000001
-----

```

Ridge for all using R2 score

```
def PolynomialRidge(train_file, test_file, header, degree):
    train_df = pd.read_csv(train_file, header=header, delimiter=",")
    test_df = pd.read_csv(test_file, header=header, delimiter=",")

    X_train = train_df.iloc[:, :-1].values
    y_train = train_df.iloc[:, -1].values

    X_test = test_df.iloc[:, :-1].values
    y_test = test_df.iloc[:, -1].values

    from sklearn.linear_model import Ridge
    from sklearn.preprocessing import PolynomialFeatures

    poly_reg=PolynomialFeatures(degree=degree)
    X_poly=poly_reg.fit_transform(X_train)

    alphas=np.array([2**i for i in range(-18, 30)])

    best_alpha, best_mse=None, -float('inf')
    for alpha in alphas:
        regressor = Ridge(alpha=alpha)
        regressor.fit(X_poly, y_train)
        y_pred = regressor.predict(poly_reg.transform(X_test))

        from sklearn.metrics import mean_squared_error,
    mean_absolute_error, r2_score
        mse = r2_score(y_test, y_pred)
        if(mse>best_mse):
            best_mse, best_alpha=mse, alpha

    return np.array([best_mse, best_alpha])

for train_dir, test_dir, header in zip(train_dirs, test_dirs,
headers):
    train_files = os.listdir(train_dir)
    test_files = os.listdir(test_dir)

    val = np.zeros(2)
    for train, test in zip(train_files, test_files):
        val += PolynomialRidge(train_dir + train, test_dir + test,
header, 2)
    print(train_dir)
    val /= len(train_files)
    val = pd.DataFrame(val, index=["R2", "Alpha"], columns=["Values"])
    print(val)
    print("-----\n")

./ele-1-5-fold/train/
Values
```



```
R2          0.733629
Alpha 56524.800002
-----
```

```
./diabetes-5-fold/train/
      Values
R2          0.361546
Alpha 419533.200002
-----
```

```
./quake-5-fold/train/
      Values
R2      5.585534e-03
Alpha 1.385431e+07
-----
```

```
./laser-5-fold/train/
      Values
R2          0.949986
Alpha 217.600002
-----
```

```
./plastic-5-fold/train/
      Values
R2          0.800643
Alpha 26252.800002
-----
```

```
for train_dir, test_dir, header in zip(train_dirs, test_dirs,
headers):
    train_files = os.listdir(train_dir)
    test_files = os.listdir(test_dir)

    val = np.zeros(2)
    for train, test in zip(train_files, test_files):
        val += PolynomialRidge(train_dir + train, test_dir + test,
header, 3)
    print(train_dir)
    val /= len(train_files)
    val = pd.DataFrame(val, index=["R2 score", "Alpha"],
columns=["Values"])
    print(val)
    print("-----\n")
```

```
./ele-1-5-fold/train/
      Values
R2 score 7.392552e-01
Alpha 1.074135e+08
-----
```

```
./diabetes-5-fold/train/
```

```
Values
```

```
R2 score    0.502298
```

```
Alpha       218.403126
```

```
-----
```

```
./quake-5-fold/train/
```

```
Values
```

```
R2 score    1.351850e-02
```

```
Alpha       1.109393e+08
```

```
-----
```

```
./laser-5-fold/train/
```

```
Values
```

```
R2 score    0.987530
```

```
Alpha       52.800002
```

```
-----
```

```
./plastic-5-fold/train/
```

```
Values
```

```
R2 score    0.814724
```

```
Alpha       25.625002
```

```
-----
```